



# CS3243: Introduction to Artificial Intelligence

Semester 2, 2019/2020

# Teaching Staff

- Lecturer: Yair Zick  
Email: [zick@comp.nus.edu.sg](mailto:zick@comp.nus.edu.sg)  
Office: COM2-02-60
- Co-lecturer: Daren Ler  
Email: [dler@comp.nus.edu.sg](mailto:dler@comp.nus.edu.sg)  
Office: COM2-02-65

# Teaching Staff

- TAs:
  - Mohit Agrawal (eo508702@u.nus.edu)
  - Hu Zikun (eo338273@u.nus.edu)
  - Lew Yin Hui Stephanie (dcslyhs@nus.edu.sg)
  - Arpan Losalka (arpan@u.nus.edu)
  - Tohar Lukov (eo384144@u.nus.edu)
- Undergraduate TAs
  - Vineeth Buddha (eo325373@u.nus.edu)
  - Li Zeyong (eo032604@u.nus.edu)
  - Teh Zi Cong Nicholas (eo148237@u.nus.edu)

# Teaching Resources: LumiNUS

<http://luminus.nus.edu.sg/>

- Lesson Plan
- Lectures, Tutorials, Supplementary Materials, Homework
- Discussion forum
  - Any questions related to the course should be raised on this forum
  - Emails to me will be considered public unless otherwise specified
- Announcements
- Homework submissions
- Webcasts

# A 'Tasting Menu' of AI

## Foundational concepts of AI

- Search
- Game playing
- Logic
- Uncertainty
- Probabilistic reasoning

## Who?

- Undergraduates
- beginning graduate students.
- CS orientation, or by permission.

# Beyond CS3243

## Machine Learning

CS3244

CS5242

CS5339

CS5340

CS5344

## Search & Planning

CS4246

CS5338,  
TBA

## Logic

CS4248

CS6207

CS4244

## ... And more!

CS4261


CS6208

CS6281

# Readings

- Textbook:

- Russell and Norvig (2010).  
Artificial Intelligence: A Modern Approach (3rd Edition ← Important!)
- Online Resources, Code, and ERRATA:  
<http://aima.cs.berkeley.edu/>
- We will not cover entire book! But it makes for an interesting read...



# Syllabus

- Introduction and Agents (chapters 1, 2)
- Search (chapters 3, 4, 5, 6)
- Logic (chapters 7, 8, 9)
- Uncertainty (chapters 13, 14)
- Learning (chapters 18, 21)



# Assessment Overview

What	When	Grade Percentage
Midterm Exam (during lecture, NO make-up)	5 March 2020	20%
Final Exam	2 May 2020 (afternoon)	50%
Class Project	Week 7 and Week 10	25%
Tutorials + Attendance	-	5%



# Introduction


AIMA Chapter 1



# What is AI?

# Making Computers (at least as) Good as Humans in Human Tasks

- ... but not necessarily in the same way a human would!
  - Birds and planes fly very differently
  - AI solves problems very differently.
  - I think!



```
graph TD; L[Linguistics] --> AI[AI]; CS[CS] --> AI; M[Math] --> AI; E[Economics] --> AI; P1[Psychology] --> AI; P2[Psychology] --> AI;
```

Linguistics

CS

Psychology

**AI**

Math

Psychology

Economics

## Philosophy

- Ethics
- Logic
- Learning
- Rationality
- Theory of the Mind

## Computer Science

- Theory of Computing
- Hardware
- Control Theory
- Dynamic Systems

## Mathematics

- Formal representation
- Probability
- Statistics

## Economics

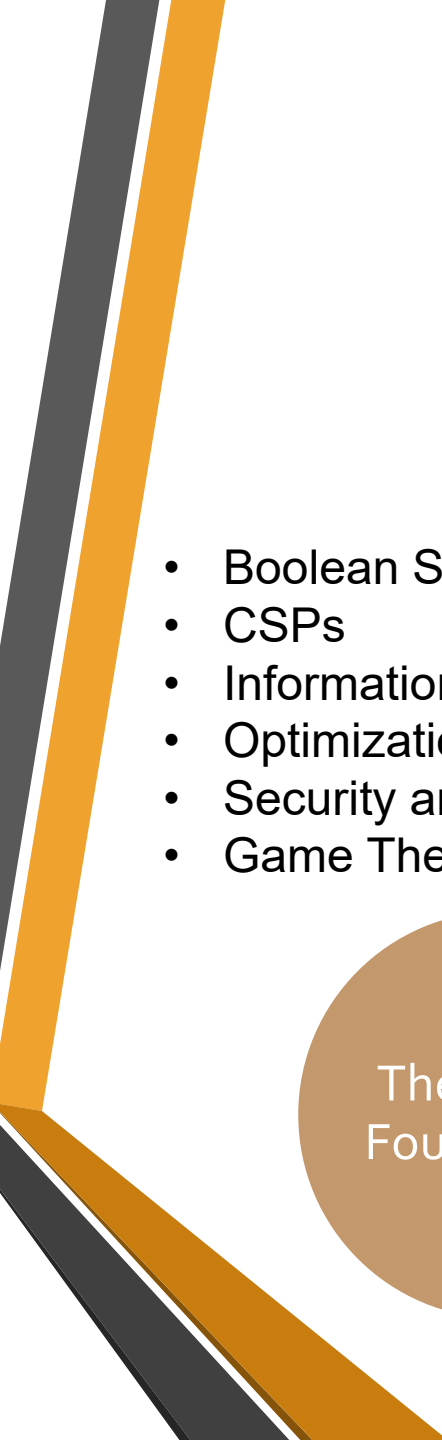
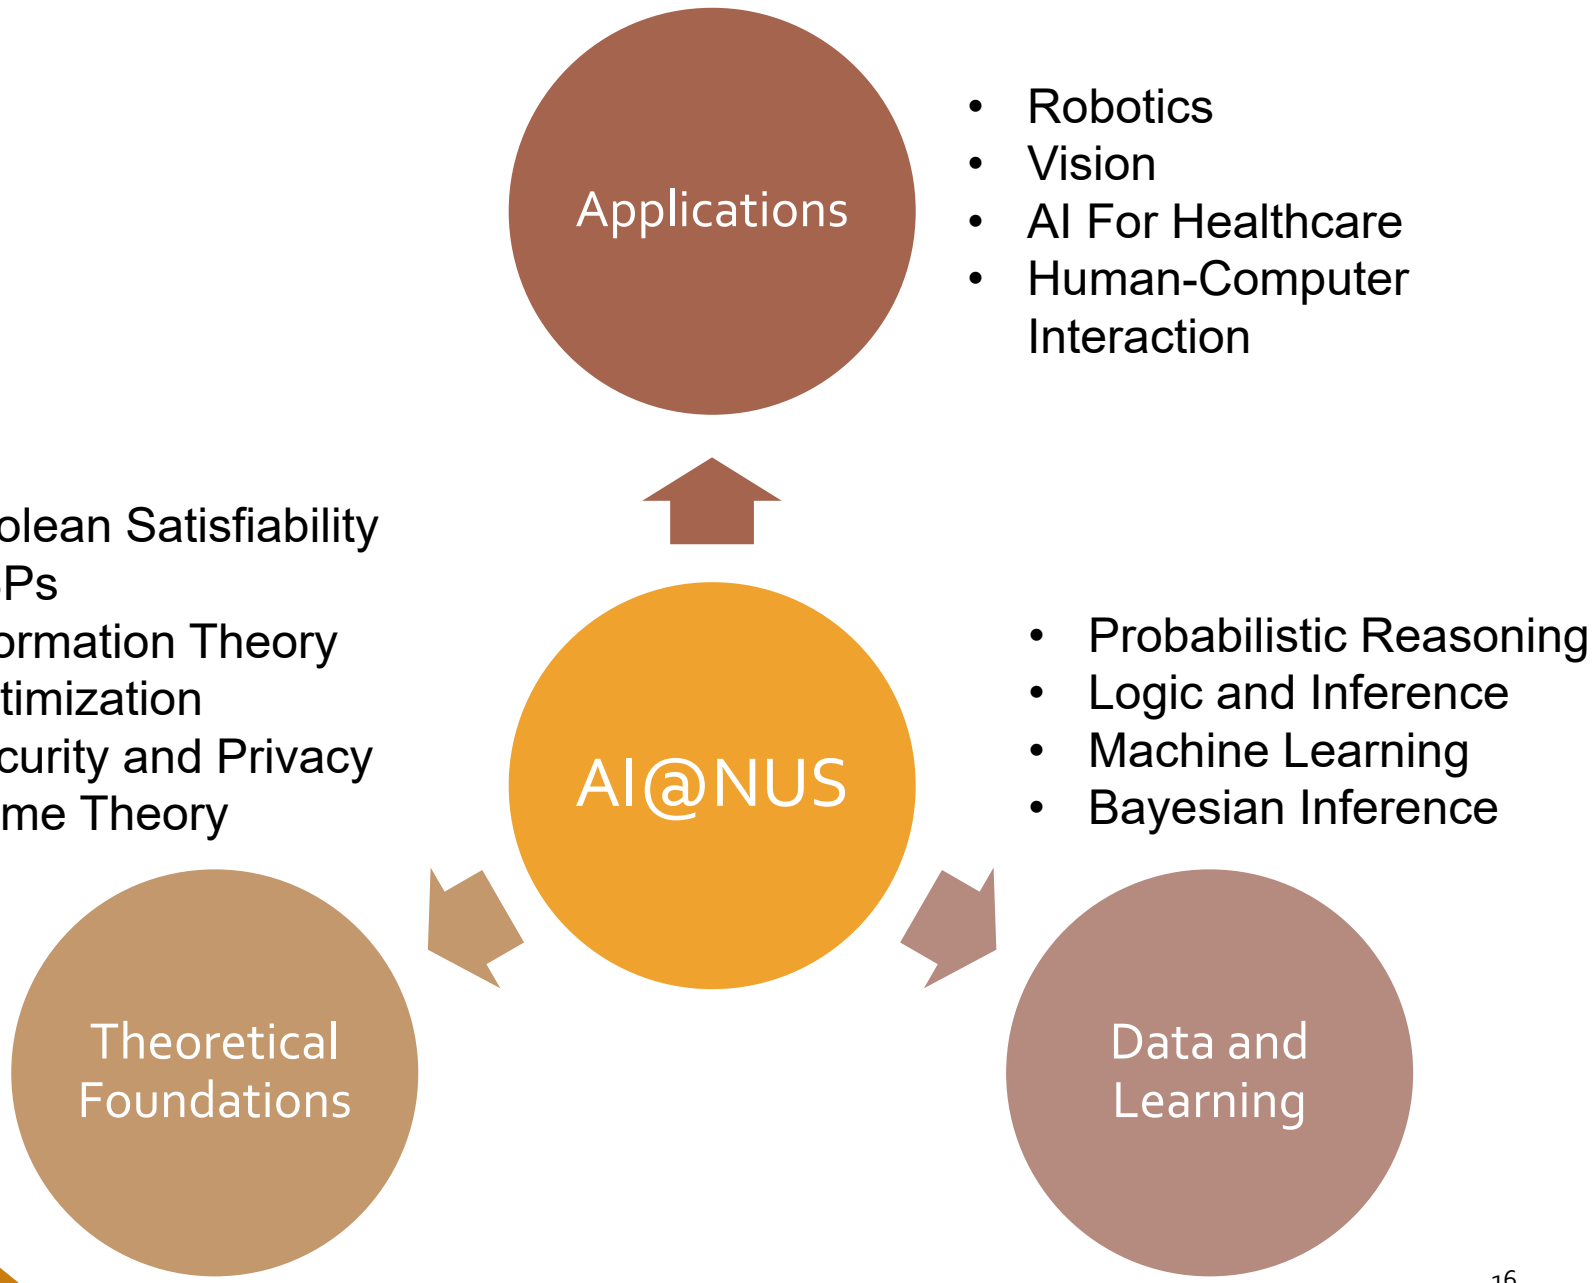
- Game theory
- Decision theory
- Fair Division
- Utility theory

## Psychology

- Perception and motor control
- Experiments

## Linguistics

- Knowledge representation
- grammar

- 
- 
- ```
graph TD; AI((AI@NUS)) --> Apps((Applications)); AI --> Theo((Theoretical Foundations)); AI --> Data((Data and Learning)); AI --> Prob((Probabilistic Reasoning));
```
- Boolean Satisfiability
  - CSPs
  - Information Theory
  - Optimization
  - Security and Privacy
  - Game Theory

Applications

- Robotics
- Vision
- AI For Healthcare
- Human-Computer Interaction

AI@NUS

Theoretical  
Foundations

Data and  
Learning

- Probabilistic Reasoning
- Logic and Inference
- Machine Learning
- Bayesian Inference




# AI is Getting Better at Gameplay

“A computer once beat me at chess, but it was no match for me in kickboxing” – Emo Philips

| Year | Game                            | Program                   | Developer             | Techniques                                 |
|------|---------------------------------|---------------------------|-----------------------|--------------------------------------------|
| 1994 | Checkers                        | Chinook                   | U. Alberta            | Rule Based + search                        |
| 1997 | Chess                           | Deep Blue                 | IBM                   | Search + randomization                     |
| 2008 | Limit Texas Hold'em             | Polaris (Cepheus 2015)    | U. Alberta            | Agent based modeling, game theory          |
| 2011 | Jeopardy                        | Watson                    | IBM                   | NLP, Information retrieval, data analytics |
| 2015 | 2 player No Limit Texas Hold'em | Claudico (later Libratus) | Carnegie Mellon Univ. | Game Theory, Reinforcement Learning        |
| 2016 | Atari Games                     | DeepMind                  | DeepMind              | Deep Learning                              |
| 2016 | Go                              | AlphaGo                   | DeepMind              | Deep Learning, search                      |
| 2018 | DOTA 2                          | OpenAI Five               | OpenAI                | RL, Deep learning, search                  |
| 2019 | Starcraft 2                     | AlphaStar                 | DeepMind              | RL, Deep learning, search                  |


# When Can a Machine Truly Think?

- Turing (1950). Computing Machinery and Intelligence: “Can machines think?” → “Can machines behave intelligently?”
- Operational test for intelligent behavior: The Imitation Game



# Is the Turing Test a Good Idea?

- What are the advantages of the Turing test?
- What are the disadvantages?



- Think about 1-2 of each
- Talk to the person next to you about it
- We'll move to entire class after

# Winograd Schema

The city councilmen refused the demonstrators a permit because they [feared/advocated] violence.

Who [feared/advocated] violence?

**Answers:** The city councilmen/the demonstrators.

# Winograd Schema

The trophy doesn't fit into the brown suitcase because it's too [small/large].

What is too [small/large]?

**Answers:**The suitcase/the trophy.

# Winograd Schema

- A sentence with an ambiguous pronoun (he/she/it/they/their ...).
- Need to decide what the pronoun refers to (offered two options).
- This depends on **context**.
- Context can be easily flipped with a **single word**.

The firemen arrived [after/before] the police because they were coming from so far away.

Who came from far away?

**Answers:** The firemen/the police.

# Winograd Schema Challenge

- You are given  $m$  Winograd schema, with the context word chosen uniformly at random.
- Design an AI that can correctly resolve a significant number of them.
- What is a trivial lower bound on the number of schema one can solve?

# A Single Test for Intelligence?

- Difficult to resolve
- Tests tend to be
  - over-specified
  - very subjective
- Result will be debatable





# Acting Rationally: Rational Agent

- **Rational** behavior: doing the “right thing”
  - What is the “right thing” to do? Expected to achieve best outcome
    - Best for whom?
    - What are we optimizing?
    - What information is available?
    - Unintended effects
- Break through wall to get a cup of coffee
  - Prescribe high doses of opiates to depressed patient
  - kill human who tries to deactivate robot

# Rational Agents

- An **agent** is an entity that perceives and acts
- **This course:** designing rational agents
- An agent is a function from **percept histories** to **actions**, i.e.,  $f: P^* \rightarrow A$
- We seek the best-performing agent for a certain task; must consider computation limits!

design best **program**  
given resources

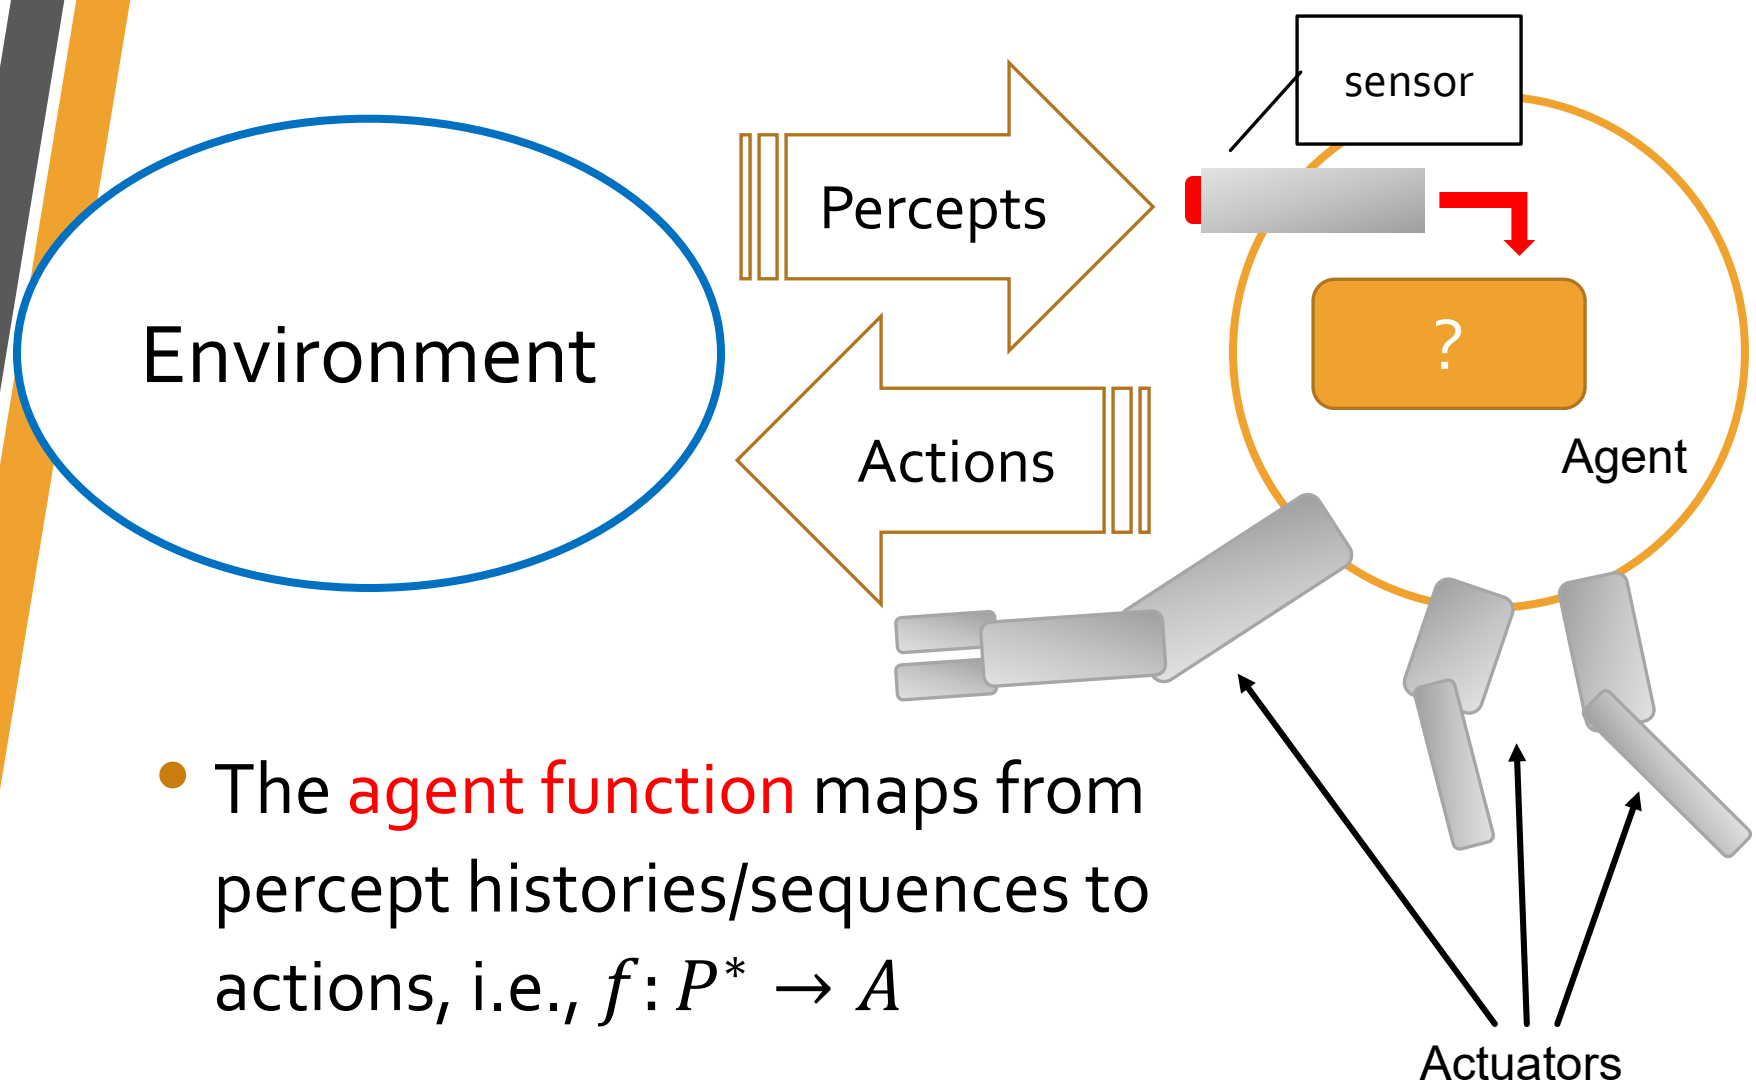


# Intelligent Agents

AIMA Chapter 2

# Agents


- Anything that can be viewed as **perceiving** its **environment** through **sensors**; **acting** upon that environment through **actuators**
- **Human agent:** eyes, ears, skin etc. are sensors; hands, legs, mouth, and other body parts are actuators
- **Robotic agent:** cameras and laser range finders for sensors; various motors for actuators



- The **agent function** maps from percept histories/sequences to actions, i.e.,  $f: P^* \rightarrow A$
- The **agent program** runs on the physical **architecture** to perform  $f$

agent = architecture + program

# Vacuum-Cleaner World



- Percepts: location and status, e.g., [*A*, Dirty]
- Actions: Left, Right, Suck, NoOp

# Vacuum-Cleaner Agent Function

## Percept Sequence

[A, Clean]

[A, Dirty]

[B, Clean]

[B, Dirty]

[A, Clean], [A, Clean]

[A, Clean], [A, Dirty]

## Action

Right

Suck

Left

Suck

Right

Suck

# Rational Agents

- An agent should strive to “do the right thing”, based on what it can perceive and the actions it can perform. The right action: maximize agent success.
- **Performance measure**: objective criterion for measuring success of an agent's behavior
- Vacuum-cleaner agent:
  - amount of dirt cleaned
  - time taken
  - electricity consumed
  - noise generated

Perhaps a bit of everything?



# Rational Agents

- **Rational Agent:**
  - For each possible percept sequence, select an action that is expected to maximize its performance measure...
  - given the evidence provided by the percept sequence and whatever built-in knowledge the agent has.

# Rational Agents

- Rationality  $\neq$  omniscience (all-knowing with infinite knowledge)
- Agents can perform actions that help them gather useful information (exploration)
- An agent is **autonomous** if its behavior is determined by its own experience (with ability to learn and adapt)

# Specifying Task Environment: PEAS

- PEAS: Performance measure, Environment, Actuators, Sensors
- Must first specify the setting for intelligent agent design
- Consider, e.g., the task of designing an automated taxi driver:
  - Performance measure
  - Environment
  - Actuators
  - Sensors



# Specifying Task Environment: PEAS

## Automated Taxi

### Performance Measure

- Safe
- Fast
- Legal
- Comfort
- Revenue

### Environment

- Roads
- Other traffic
- Pedestrians
- Customers

### Actuators

- Steering wheel
- Accelerator
- Brake
- Signal
- Horn

### Sensors

- Camera
- Sonar
- Speedometer
- GPS
- Engine sensors

# Specifying Task Environment: PEAS

## Part Picking Robot

### Performance Measure

- % parts in correct bins

### Environment


- Conveyor belt
- parts
- bins

### Actuators

- Jointed arm
- hand

### Sensors

- Camera
- joint angle sensors



# Specifying Task Environment: PEAS

## Medical Diagnosis System

### Performance measure

- Healthy patient
- cost
- lawsuits

### Environment

- Patient
- hospital
- staff

### Actuators

- Screen display (questions, tests, diagnoses, treatments, referrals)

### Sensors

- Keyboard
- Medical Readings
- Medical History



# Specifying Task Environment: PEAS


## Interactive English Tutor

### Performance measure

- Student's score on test

### Environment

- Set of students
- Testing agency
- Chat platform



### Actuators

- Screen display (exercises, suggestions, corrections)

### Sensors

- Keyboard entry

# Properties of Task Environments

## Fully observable (vs. partially observable):

- sensors provide access to the complete state of the environment at each point in time.

## Deterministic (vs. stochastic)

- The next state of the environment is completely determined by the current state and the action executed by the agent.

## Episodic (vs. sequential)

- The choice of **current** action does not depend on actions in past episodes.
- Alternatively: **current** action does not affect future decisions



# Properties of Task Environments

## Static (vs. dynamic)

- The environment is unchanged while an agent is deliberating.

## Discrete (vs. continuous)

- A finite number of distinct states, percepts, and actions.

## Single agent (vs. multi-agent)

- An agent operating by itself in an environment.

# Properties of Task Environments

| Task Environment | Crossword puzzle | Part-picking robot | Taxi driving |
|------------------|------------------|--------------------|--------------|
| Fully observable | Yes              | No                 | No           |
| Deterministic    | Yes              | No                 | No           |
| Episodic         | No               | Yes                | No           |
| Static           | Yes              | No                 | No           |
| Discrete         | Yes              | No                 | No           |
| Single agent     | Yes              | Yes                | No           |

Properties of task environment largely determine agent design. World is partially observable, stochastic, sequential, dynamic, continuous, multi-agent.

# Agent Functions and Programs

- An agent is completely specified by the agent function mapping percept sequences to actions
- One agent function (or a small equivalence class) is rational
- Aim: Find a way to implement the rational agent function concisely

# Agent Functions and Programs

**Solution:** write all possible percepts and optimal actions in a table.

All done!

# Table-Lookup Agent

```
function TABLE-DRIVEN-AGENT(percept) returns action  
  static: percepts, a sequence, initially empty  
           table, a table of actions, indexed by percept sequences, fully specified  
  
  append percept to the end of percepts  
  action  $\leftarrow$  LOOKUP(percepts, table)  
  return action
```

- Drawbacks:
  - Huge table to store
  - Take a long time to build the table
  - No autonomy: impossible to learn all correct table entries from experience
  - No guidance on filling in the correct table entries

# Agent Types

- Four basic types in order of increasing generality:
  - Simple reflex agent
  - Model-based reflex agent
  - Goal-based agent
  - Utility-based agent

# Reflex Agent

- Passive: only acts when observes a percept
- Updates *state* based on *percept* only.
- Easy to implement

Motion sensor triggered.

*state*  $\leftarrow$  *WELCOME MODE*



# Model Based Reflex Agent

- Passive: only acts when observes a percept
- *state* is updated based on **percept**, current *state*, most recent *action*, and **model of world**.

Motion sensor triggered. I am on ACTIVE  
MODE. I AUTHORIZED ACCESS.  $\Rightarrow$   
HUMAN IN HOUSE

*state*  $\leftarrow$  WELCOME MODE






# Goal Based Agent

- Has **goals**, acts to achieve them (not passive).
- *state* is updated based on **percept**, current *state*, most recent *action*, and **model of world**.

Motion sensor triggered. I am on **ACTIVE MODE**. I **AUTHORIZED ACCESS**.  $\Rightarrow$  **HUMAN IN HOUSE**

*state*  $\leftarrow$  *WELCOME MODE*(1)

Goal: make human happy  
(and buy products from my  
vendor)




# Utility Based Agent

- Has **utility function**, acts to maximize it.
- *state* is updated based on **percept**, current *state*, most recent *action*, and **model of world**.

$$\max \alpha \times \text{happiness} + \beta \times \text{purchases}$$

Motion sensor triggered. I am on **ACTIVE MODE**. I **AUTHORIZED ACCESS**.  $\Rightarrow$  **HUMAN IN HOUSE**

$state \leftarrow WELCOME\ MODE(17)$



# Learning Agent

Performance  
Standard

17 sodas is  
unacceptable!

Human did not  
like 17 sodas ☹️

Update  $\alpha, \beta$   
configuration

Critic

Learner

Problem  
Generator

Performance  
Element

(think the utility agent architecture  
from the last slide)

Maybe 170 sodas  
this time?



# Exploitation vs. Exploration

- An agent operating in the real world must often choose between:
  - maximizing its expected utility according to its current knowledge about the world; and
  - trying to learn more about the world since this may improve its future gains.

***Exploitation*** vs. ***Exploration***